



UNIVERSITÀ  
DEGLI STUDI DI BARI  
ALDO MORO

# Amazon Electronics Reviews

A Big Data Pipeline for Business Intelligence



**Vito Marco Rubino**

MSc in *Data Science – Big Data Management & Analytics Project*  
**Academic Year 2025 – 2026**

# Table of contents

**01**

Introduction

**02**

Dataset

**03**

Design

**04**


Modeling

**05**

Results

**06**

Conclusions



01

# Introduction



# Introduction


Modern e-commerce environments generate massive volumes of heterogeneous data

## Challenges

- **Schema rigidity**  
traditional RDBMS cannot efficiently handle the schema variability of products catalogue
- **Scalability**  
Processing million of data exceed vertical scaling capacities

## Goals

1. Build an **ETL pipeline** to handle semi-structured documents
2. Apply **NLP** to extract latent drivers of customer satisfaction
3. Use **PageRank** to find network influencers and **ALS** for personalization




02

# Dataset

# The Amazon Review dataset

This study utilizes the **Amazon Review Dataset**, specifically focusing on the **Electronics** category subset.



**20,994,353**

The number of **reviews**

**786,868**

The number of **products** metadata

# Data Ingestion & Cleaning

Modern e-commerce generates massive mixed-format data, with challenges including  
**type inconsistencies, HTML tags and Python literals**



A two stage pipeline was implemented to address these challenges

1

## Data Ingestion

Applied Safe Mapping to parse literals into standard JSON.



Parsing strategy

```
1 # Context mapping for safe evaluation
2 safe_context = { "null": None, "nan": float('nan'), "true": True }
3 # Yield valid JSON string for Spark ingestion
4 yield json.dumps(eval(line, safe_context))
```

2


## Data Cleaning

Removed HTML tags, cast timestamps, and handled missing values using Spark SQL



Data Cleaning

```
1 // HTML Removal & Date Conversion
2 .withColumn("reviewText", regexp_replace($"reviewText", "<[^>]+>", " "))
3 .withColumn("reviewDate", to_date(from_unixtime($"unixReviewTime")))
4 // Handling Nulls in Votes & Key Standardization
5 .withColumn("helpful_votes", coalesce($"vote".cast(IntegerType), lit(0)))
6 .withColumnRenamed("asin", "productID")
```



03

Design

# Schema Design


Data is modeled using a **star-schema**

## Fact (*Reviews*)

stores only quantitative measures (ratings, votes) and unstructured text.

## Dimensions (*Products* and *Users*)

Describe the fact providing the context for users and products metadata




# Optimization strategy

To answer business questions reviews must be linked with the descriptive dimension tables


## Bottleneck

moving 20M+ reviews for a  
Sort-Merge join triggers a  
Full Network Shuffle



## Broadcast Join

small table broadcasting enables local  
Map-Side Joins, joining *Reviews*  
in-memory without network shuffle



# 04

## Modeling

# Pipeline Overview

The analytical workflow unfolds in three distinct modeling phases

01.

## Natural Language Processing

Extracting sentiment drivers and semantic context using also **Word2Vec**

02.

## Network Analysis

Applying **PageRank** algorithm to identify market influencers

03.

## Predictive Modeling

Forecasting user preferences using **ALS** collaborative filtering

# Sentiment Analysis and Drivers Extraction

A supervised classification pipeline quantify the specific vocabulary driving customer satisfaction



## Setup

Binary labeling of ratings ( $>3.0$ ) and 80-20 train-test splitting



## NLP pipeline

```
1 val data = cleanReviews.withColumn("label",
2   when($"overall" > 3.0, 1.0).otherwise(0.0))
3
4 val Array(trainingData, testData) = data.randomSplit(Array(0.8, 0.2), seed =
42)
```



## Feature engineering

Custom stop-word removal, rare term filtering, and IDF weighting for common words



```
1 // Stop-words removal
2 val remover = new StopWordsRemover().setStopWords(defaultStopWords ++
customArray)
3
4 // Produces a sparse-vector with term-frequency and location
5 val cv = new CountVectorizer().setVocabSize(10000).setMinDF(10.0)
6
7 val idf = new IDF().setInputCol("rawFeatures").setOutputCol("features")
```



## Drivers Extraction


Logistic Regression training and coefficients ranking



```
2 val lr = new LogisticRegression().setMaxIter(10).setRegParam(0.01)
3 // [...]
4 val model = pipeline.fit(trainingData)
5 val predictions = model.transform(testData)
6 // [...]
7 val wordWeights = cvModel.vocabulary.zip(lrModel.coefficients.toArray)
8
9 println("Top 10 Positive Words:")
10 wordWeights.sortBy(-_._2).take(10).foreach(println)
```

# Semantic modeling via Word2Vec

Word2Vec embeds words into a space where proximity encodes semantic similarity enabling semantic disambiguation and topic discovery through word clustering



## Model Configuration

- **Vocabulary ( $V$ )** derived from `filtered_words`, using `minCount = 50` to remove noise (typos)
- **Feature space ( $N = 100$ )** compresses  $V$  into a dense feature space
- **Context Scope** optimize weight with a 5-token sliding window

```
1 val word2VecEstimator = new Word2Vec()
2   .setInputCol("filtered_words")
3   .setOutputCol("result_vector")
4   .setVectorSize(100) // Dimensionality of the feature space
5   .setMinCount(50)   // Minimum frequency to include a token
6   .setWindowSize(5)  // Context window size
```

# Graph-based Network Analysis

The products network was modeled as a direct graph to quantify structural influence following three steps



## Graph Analysis

```
1 // Map String IDs to Long using hashCode
2 val verticesRDD: RDD[(Long, String)] = productsDF.rdd.map(row => {
3   (row.getString(0).hashCode.toLong, row.getString(1))
4 })
5
6 // Edge Generation: Explode 'also_buy' array to create directed links
7 val edgesRDD: RDD[Edge[Int]] = distinctEdgesDF.rdd.map(row => {
8   val sourceId = row.getString(0).hashCode.toLong
9   val targetId = row.getString(1).hashCode.toLong
10  Edge(sourceId, targetId, 1) // Unweighted edge
11 })
12
13 val productGraph = Graph(verticesRDD, edgesRDD)
14
15 // Execute PageRank
16 val pageRankGraph = productGraph.pageRank(0.001)
```



## Vertex mapping

Conversion of alphanumeric ASINs into long integers via hashing



## Edges generation

Explosion of **also\_buy** arrays into directed edge (Source → Target)



## PageRank

Identify network structural hubs

# Recommender System via ALS

Implementation of a Collaborative Filtering using ALS to predict latent user preferences

01.

## Input

We start with a *Reviews* ( $R$ ) matrix that is at most empty

	Headset	Mouse	Laptop
Alice	5	3	?
Bob	4	?	?
Charlie	1	1	5

02.

## Factorization

ALS breaks  $R$  into two smaller and dense matrices indicating User and Item factors

$$R \approx U \times V^T$$

03.


## Optimization

solve for one matrix while fixing the other using Least Squares until convergence

04.

## Prediction



Predict a missing rating for User  $i$  and item  $j$   
 $\hat{r}_{ij} = u_i \cdot v_j^T$



# 05

## Results

# Global Sentiment Drivers



Logistic Regression coefficients reveal a **clear split** in consumer drivers:




## Positive predictors

Fueled by emotional satisfaction and user experience




## Negative predictors

Driven by reliability and durability issues (e.g. "stopped", "returned")



# Synonym Discovery

**Word2Vec** mapped synonyms into a continuous vector space



High similarity scores confirm it learned synonyms:

- **Hardware:** “*display-screen*” (0.89)
- **Failure:** “*defective-broken*” (0.88)

‘screen’		‘broken’	
Associated Term	Similarity	Associated Term	Similarity
display	0.89	defective	0.88
lcd	0.85	dead	0.81
monitor	0.81	stopped	0.79
pixels	0.76	cracked	0.75

# Graph Network Centrality

PageRank identifies **low-cost accessories** as the primary structural hubs

The ranking reveals a topology hierarchy:


- **Connectors dominance**

Accessories serve as primary hubs and bridge product clusters

- **Flagship Endpoints**

High-value devices (TVs, cameras) act as graph endpoints rather than hubs

Rank Score	Product Name
524.30	AmazonBasics High-Speed HDMI Cable
412.15	SanDisk Ultra 32GB MicroSDHC
380.50	Panasonic BK-3MCCA8 Eneloop AA Batteries
310.20	Belkin Surge Protector
295.00	Chromecast Google



# Recommender System via ALS

## Quantitative Evaluation

- ALS got a RMSE of 1.12 on a 5-star scale
- The model's prediction deviates, on average, by roughly 1 star from the actual rating

## Qualitative Evaluation on specific user 'A5VBZG7I2RL1'

- Classified as a "Peripheral User" based on previous reviews (e.g., Bluetooth dongle).
- Top 3 suggestions target peripherals and storage expansion.

Product Title	Price (\$)	Predicted Rating
Logitech M510 Wireless Mouse	19.99	4.85
Anker 4-Port USB 3.0 Hub	9.99	4.72
WD 2TB Elements Portable Drive	69.00	4.65

# Market Intelligence

Spark SQL aggregations reveal strategic positioning and quality alerts.


## Revenue Strategy

- **Premium leaders:** Apple, Samsung
- **Volume drivers:** accessories by Belkin and Logitech



## Quality Control

**Audio & Video Accessories** has 45% of reviews being **negative** (rating < 3.0)



# Product and Community

Drilling down to identify Key Opinion Leaders and extract specific product keywords


Querying user activity reveals:

- **Top influencers** by reviews count and helpful votes
- **Lon J. Seidman** holds the highest review count in the Electronics category

Reviewer Name	User ID	Reviews	Helpful Votes	Avg Score
Lon J. Seidman	A26877IWJGISYM	21	3030	4.24
Mark	A17BUUBOU0598B	17	1066	3.65
NLee the Engineer	AOEAD7DPLZE53	12	804	4.08
M. JEFFREY MCMAHON	A3NCIN6TNLOMGA	12	755	4.58
Darren Levine	A1UIMU8Q87ZPCH	12	697	4.33

Category	Distinctive Keywords
Pros	range, setup, outdoor, christmas, lights
Cons	battery, stopped, months, died, plastic


- **Split** data into positive ( $\text{rating} > 3$ ) and negative ( $\text{rating} \leq 3$ ) and extract frequent words
- Extract keywords applying **Differential Set Logic**
$$\text{Words}_+ = \text{Words}_+ - (\text{Words}_+ \cap \text{Words}_-)$$
- Specific Product Analysis for Woods Remote identified low-temperature failure ("cold", "died")



# 06



# Conclusions



# Key takeaways

## Scalability

MongoDB and Spark to process 20M+ records via a distributed architecture

## Semantic Reasoning

Word2Vec and Logistic Regression decoded the language of the market

## Network Insights


We proved low-cost accessories are the true hubs instead of flagship products

## Collaborative Filtering

ALS powered the generation of personalized recommendation

What if we could evolve these models through real-time incremental learning?

Thank you!



   
**Vito Marco Rubino**  
MSc in *Data Science*  
**A.Y. 2025 - 2026**